

Améliorer la qualité numérique des calculs

Licence informatique (L3)

Ph. Langlois

Université de Perpignan, France

4 avril 2018



Context and motivations

Sources of errors when computing the solution of a scientific problem in floating point arithmetic:

- mathematical model,
- truncation errors,
- data uncertainties,
- **rounding errors.**

Context and motivations

Sources of errors when computing the solution of a scientific problem in floating point arithmetic:

- mathematical model,
- truncation errors,
- data uncertainties,
- **rounding errors.**


Rounding errors may totally corrupt a floating point computation:

- accumulation of billions of floating point operations,
- intrinsic difficulty to solve the problem accurately.

Example: polynomial evaluation

Evaluation of univariate polynomials with floating point coefficients:

- the evaluation of a polynomial suffers from rounding errors
- example : in the neighborhood of a multiple root




$p(x) = (x - 2)^9$ in expanded form
near the multiple root $x = 2$

Example: polynomial evaluation

Evaluation of univariate polynomials with floating point coefficients:

- the evaluation of a polynomial suffers from rounding errors
- example : in the neighborhood of a multiple root



$p(x) = (x - 2)^9$ in expanded form
near the multiple root $x = 2$
evaluated with the Horner algorithm
in IEEE double precision.

General motivation

How to:

- improve and validate the accuracy of a floating point computation,
- recover the numerical reproducibility of parallel floating-point computation,
- without large computing time overheads ?

- More hardware precision
- Software simulation of more computing precision
- More accurate algorithms
- As fast as possible accurate algorithms
- Parallelism is everywhere
- Reproducibility at least to debug!

Main issues today

Starting point: IEEE-754 floating point arithmetic

- Best possible accuracy for $+$, $-$, \times , $/$, $\sqrt{}$
- Add is **not associative**

Summing n floating numbers : focus on accuracy

- Core computation, numerous algos, recently some really smart ones
- Computed sum accuracy : doubled or more, faithful or correctly rounded

Summing n floating numbers : focus on running-time and memory print

- Running-time and memory print are discriminant factors when computing the best possible accurate sum
- Appreciating the actual performance of one algo is not an easy task :
flop/s ? hardware counters ? compiler options ?

Reliable and significant measure of the time complexity?

The classic way: count the number of flop

- A usual problem: double the accuracy of a computed result
- A usual answer for polynomial evaluation (degree n)

Metric	Horner	CompHorner	DDHorner
Flop count	$2n$	$22n + 5$	$28n + 4$
Flop count ratio	1	≈ 11	≈ 14
Measured #cycles ratio	1	$2.8 - 3.2$	$8.7 - 9.7$

Flop count vs. run-time measures

- Flop counts and measured run-times are not proportional
- Run-time measure is a **very** difficult experimental process
- Which one trust?

Reproducible HPC numerical simulations?

Numerical reproducibility of parallel computation

- Getting bitwise **identical** results for every p -parallel run, $p \geq 1$

One industrial scale simulation code

- Simulation of free-surface flows in 1D-2D-3D hydrodynamics
- 300 000 loc. of open source Fortran 90
- 20 years, 4000 registered users, EDF R&D + international consortium


Telemac 2D [5]

- 2D hydrodynamic: Saint Venant equations
- Finite element method, triangular element mesh, sub-domain decomposition for parallel resolution
- Mesh node unknowns: water depth (H) and velocity (U,V)





A complex Telemac 2D simulation

The Malpasset dam break (1959)


- A five year old dam break: 433 dead people and huge damage
- Triangular mesh: 26000 elements and 53000 nodes
- Simulation: →35min. after break with a 2sec. time step



Profondeur d'eau obtenue pour t=2200s



How to trust this complex simulation?



A reproducible simulation?

	velocity U	velocity V	depth H
The sequential run	0.4029747E-02	0.7570773E-02	0.3500122E-01
one 64 procs run	0.4935279E-02	0.3422730E-02	0.2748817E-01
one 128 procs run	0.4512116E-02	0.7545233E-02	0.1327634E-01

Reproducibility failure

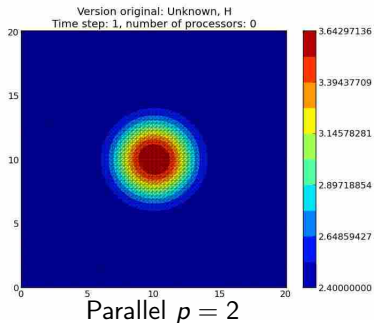

- Up to $\times 2.5$ uncertainty
- The privileged sequential run?

Telemac2D: the simplest goutedo simulation

The goutedo simulation test case

- 2D-simulation of a water drop fall in a square basin
- Unknown: water depth for a 0.2 sec time step
- Triangular mesh: 8978 elements and 4624 nodes



Expected numerical reproducibility (time step = 1, 2, ...)



A white plot displays a non-reproducible value

Numerical reproducibility?


time step = 1




A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 2



Sequential




Parallel $p = 2$

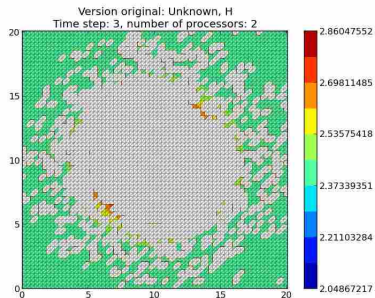
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 3



Sequential




Parallel $p = 2$

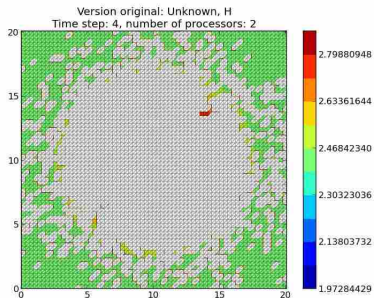
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 4



Sequential




Parallel $p = 2$

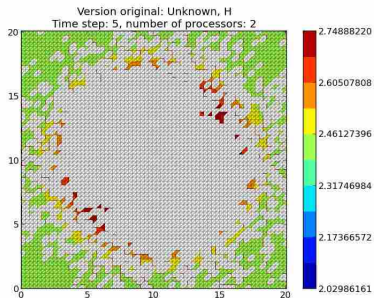
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 5



Sequential




Parallel $p = 2$

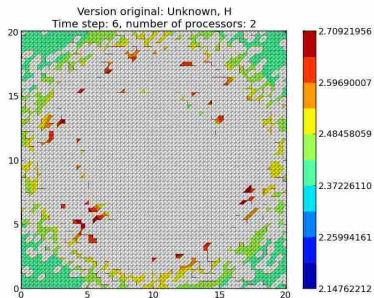
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 6



Sequential




Parallel $p = 2$

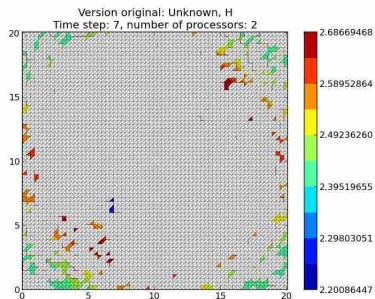
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 7



Sequential




Parallel $p = 2$

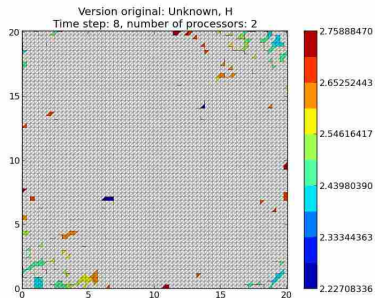
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 8



Sequential




Parallel $p = 2$

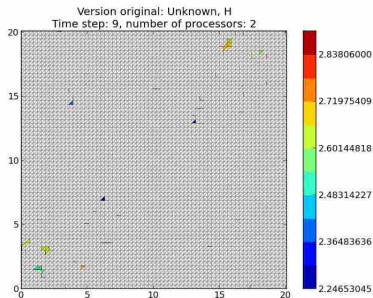
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 9



Sequential




Parallel $p = 2$

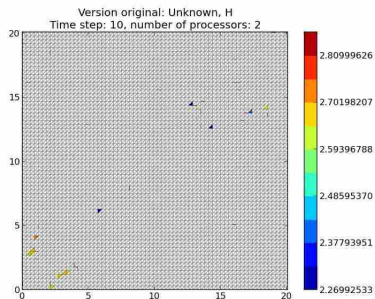
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 10



Sequential




Parallel $p = 2$


A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 11



Sequential




Parallel $p = 2$


A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 12



Sequential




Parallel $p = 2$

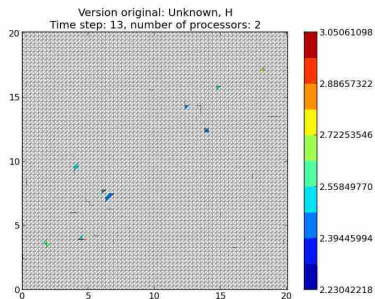
A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 13



Sequential




Parallel $p = 2$


A white plot displays a non-reproducible value

Numerical reproducibility?

time step = 14



Sequential




Parallel $p = 2$


A white plot displays a non-reproducible value

NO numerical reproducibility!

time step = 15




Sequential




Parallel $p = 2$

NO numerical reproducibility!



Sequential



Parallel $p = 2$

Part 1: More accuracy

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms

- Some more results

4 More accuracy: conclusion

Part 2: More reproducibility

Part 3: More performance

Part I

More accuracy

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms

- Some more results

4 More accuracy: conclusion

How to manage accuracy and speed?

So many ways . . . and a new “best one” every year since 1999

1965 Møller, Ross

1969 Babuska, Knuth

1970 Nickel

1971 Dekker, Malcolm

1972 Kahan, Pichat

1974 Neumaier

1975 Kulisch/Bohlender

1977 Bohlender, Mosteller/Tukey

1981 Linnaïmaa

1982 Leuprecht/Oberaigner

1983 Jankowski/Semoktunowicz/-
Wozniakowski

1985 Jankowski/Wozniakowski

1987 Kahan

1991 Priest

1992 Clarkson, Priest

1993 Higham

1997 Shewchuk

1999 Anderson

2001 Hlavacs/Uberhuber

2002 Li et al. (XBLAS)

2003 Demmel/Hida, Nievergelt,
Zielke/Drygalla

2005 Ogita/Rump/Oishi,
Zhu/Yong/Zeng

2006 Zhu/Hayes

2008 Rump/Ogita/Oishi


2009 Rump, Zhu/Hayes

2010 Zhu/Hayes


IEEE-754 floating point arithmetic

Floating-point numbers (normal, non zero)

$$x = (-1)^s \cdot m \cdot 2^e = \pm \underbrace{1.x_1x_2 \dots x_{p-1}}_{p \text{ bits of mantissa}} \times 2^e,$$



Rounding



- The standard model:

$$\text{fl}(a \circ b) = (1 + \varepsilon)(a \circ b), \text{ with } |\varepsilon| \leq \mathbf{u} = 2^{-p}, \text{ or } 2^{1-p}.$$

IEEE-754 (1985, 2008)

- formats, rounding modes : $+$, $-$, \times , $/$, $\sqrt{}$ are as accurate as possible, exceptions
- $\mathbf{u} = 2^{-53} \approx 10^{-16}$ for b64 in IEEE-754

Using the standard model

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms


- Some more results

4 More accuracy: conclusion

Floating-point summation: classic results [HFPA,09],[ASNA,06]

Accuracy for backward stable algorithms

- Accuracy of the computed sum $\leq (n - 1) \times \text{cond} \times u$
- $\text{cond}(\sum x_i) = \sum |x_i| / |\sum x_i|$
- No more significant digit in IEEE-b64 for large cond, i.e. $> 10^{16}$
- The length also matters for large n



Backward stable summation algorithms

Algorithm (\mathcal{A} describes a class of summation algorithms)

Let $S = \{x_1, x_2, \dots, x_n\}$.

while S contains more than one element do

 Remove two numbers x and y from S and add $x \oplus y$ to S ;

end while

Return the remaining element of S .

- Recursive summation...
- ... with increasing order sorting (IOS): $|x_1| \leq |x_2| \leq \dots \leq |x_n|$,
- ... with decreasing order (DOS),
- insertion summation with IOS,
- pairwise summation.

Accuracy of the computed sum $\leq (n - 1) \sum |x_i| \mathbf{u}$




Proposition (Another bound)

$|\hat{s}_n - s_n| \leq \mathbf{u} \sum_{i=1}^{n-1} |T_i|$, where T_i is the i -th partial sum in \mathcal{A} .

- Minimize this bound or at least every $|T_i|$.
- All the x_i have the same sign: recursive with IOS (good), insertion with IOS (the best).
- Cancellation when $\sum |x_i| \gg |\sum x_i|$: DOS better than IOS.

Conclusion of the basic steps

It exists no universally better accurate summation algorithm



Computing sums more accurately: step 1

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms

- Some more results


4 More accuracy: conclusion

More accuracy but still conditioning: overview

Arbitrary precision but *accuracy* still condition dependant

- Simulating more *precision*: double-double, quad-double, \dots , MPFR
- Compensated algorithms: Kahan(65), \dots , Sum2(05), SumK(05)
- Accuracy $\lesssim \mathbf{u} + \text{cond} \times \mathbf{u}^K$
- No more n dependency while $n\mathbf{u} \leq 1$.

T. OGITA, S. M. RUMP, AND S. OISHI




More accuracy but still conditioning: overview

Arbitrary precision but *accuracy* still condition dependant

- Simulating more *precision*: double-double, quad-double, \dots , MPFR
- Compensated algorithms: Kahan(65), \dots , Sum2(05), SumK(05)
- Accuracy $\lesssim \mathbf{u} + \text{cond} \times \mathbf{u}^K$
- No more n dependency while $n\mathbf{u} \leq 1$.

T. OGITA, S. M. RUMP, AND S. OISHI



One old and famous solution

Algorithm (Kahan's compensated summation (1965))

```
s = x1;  
c = 0;  
for i = 2 : n do  
    y = xi ⊖ c;  
    t = s ⊕ y;  
    c = (t ⊖ s) ⊖ y;  
    s = t  
end for
```


- Accuracy improvement in practice: $-c$ approximates $s \oplus y$'s rounding error.
- Compensated sum accuracy $\leq 2\mathbf{u} + n \sum |x_i| \mathcal{O}(\mathbf{u}^2)$.

One example

IEEE double precision numbers: $x_1 = 2^{53} - 1$, $x_2 = 2^{53}$ and $x_3 = -(2^{54} - 2)$.

Exact sum: $x_1 + x_2 + x_3 = 1$.

Classic summation




Relative error = 1

One example

IEEE double precision numbers: $x_1 = 2^{53} - 1$, $x_2 = 2^{53}$ and $x_3 = -(2^{54} - 2)$.


Exact sum: $x_1 + x_2 + x_3 = 1$.

Classic summation



Relative error = 1

Compensation of the rounding errors




The exact result is computed

One example

IEEE double precision numbers: $x_1 = 2^{53} - 1$, $x_2 = 2^{53}$ and $x_3 = -(2^{54} - 2)$.


Exact sum: $x_1 + x_2 + x_3 = 1$.

Classic summation



Relative error = 1

Compensation of the rounding errors



The exact result is computed

The rounding errors are computed thanks to *error-free transformations*.

Compensation

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms


- Some more results

4 More accuracy: conclusion

The underlined idea of compensated algorithm

Compensation

- 1 Let's compute the generated rounding errors
- 2 and use it **further** to compensate the whole computation.



Compensating: the later the better?

Approximations in Kahan's compensated sum

- $x_1 = 2^{p+1}, x_2 = 2^{p+1} - 2, x_3 = x_4 = \dots = x_6 = -(2^p - 1)$
computed sum = 1, compensated sum = 3, exact sum = 2 (TP).
- $|s| \geq |y|$ or “exponent-wise” at least, round-to-nearest, order 2 rounding error in y .

Pichat and Neumaier's compensated summation (72,74)

- Accumulate rounding errors and one final correction.
- Compensated accuracy $\leq \mathbf{u} |\sum x_i| + (0.75n^2 + n)\mathbf{u}^2 \sum |x_i|$.
- No more cond, no initial sort and $n \leq n'_{max} = 1/3\mathbf{u}$.

Priest's doubly compensated summation (92)

- Doubly compensated accuracy $\leq 2\mathbf{u} |\sum x_i|$.
- No more cond but initial sort of the summands and $n \leq n_{max} = 1/4\mathbf{u}$.

Doubly compensated summation (92)

Algorithm (Priest (1992))

- Sort the x_i so that $|x_1| \geq |x_2| \geq \dots \geq |x_n|$

$s_1 = x_1; c_1 = 0$

for $k = 2 : n$ do

$y_k = c_{k-1} + x_k;$

$u_k = x_k - (y_k - c_{k-1});$

$t_k = y_k + s_{k-1};$

$v_k = y_k - (t_k - s_{k-1});$

$z_k = u_k + v_k;$

$s_k = t_k + z_k;$

$c_k = z_k - (s_k - t_k)$

end for

- Need to sort the entries (need to know them!)
- Doubly compensated sum: almost correctly rounded sum!
accuracy $\leq 2u \left| \sum x_i \right|$ when $n \leq 2^{p-3}$ (precision p).

Computing sums more accurately: step 2

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further


3 More EFT and accurate algorithms

- Some more results

4 More accuracy: conclusion

Next step: towards the best accuracy

Distillation: iterate until faithful or correct rounding



- Error free transformation (EFT) $[x] \rightarrow [x^{(1)}] \rightarrow \dots \rightarrow [x^*]$
s.t. $\sum x_i = \sum x_i^*$ et $[x^*]$ returns the expected rounded value.
- Kahan (87), ..., Zhu-Hayes: **iFastSum** (SISC-09)

More space to keep everything


- Long accumulator, hardware oriented: Malcolm (71), Kulish (80)
- Cutting the summands: **AccSum** (SISC-08), **FastAccSum** (SISC-09)
- Summation at a given exponent: **HybridSum** (SISC-09), **OnLineExact** (TOMS-10)

From faithful rounding to correct rounding

- Choosing the “right side”: expensive in the break-point neighborhood
- e.g. $1 + 2^{-53} \pm 2^{-106}$

Next step: towards the best accuracy

Distillation: iterate until faithful or correct rounding



- Error free transformation (EFT) $[x] \rightarrow [x^{(1)}] \rightarrow \dots \rightarrow [x^*]$
s.t. $\sum x_i = \sum x_i^*$ et $[x^*]$ returns the expected rounded value.
- Kahan (87), ..., Zhu-Hayes: **iFastSum** (SISC-09)

More space to keep everything

- Long accumulator, hardware oriented: Malcolm (71), Kulish (80)
- Cutting the summands: **AccSum** (SISC-08), **FastAccSum** (SISC-09)
- Summation at a given exponent: **HybridSum** (SISC-09), **OnLineExact** (TOMS-10)


From faithful rounding to correct rounding

→ **Running-time and memory print are the discriminant factors**

EFT to sum two floating-point numbers

2Sum (Knuth, 65), Fast2Sum (Dekker, 71) for base ≤ 2 and RTN.

$$a + b = x + y, \text{ with } a, b, x, y \in \mathbb{F} \text{ and } x = a \oplus b.$$



Algorithm (Knuth)

```
function [x,y] = 2Sum(a,b)
    x = a  $\oplus$  b
    z = x  $\ominus$  a
    y = (a  $\ominus$  (x  $\ominus$  z))  $\oplus$  (b  $\ominus$  z)
```


Algorithm ($|a| > |b|$, Dekker)

```
function [x,y] = Fast2Sum(a,b)
    x = a  $\oplus$  b
    z = x  $\ominus$  a
    y = b  $\ominus$  z
```

Distillate $\sum x_i, x_i \in \mathbb{F}$

x_1	x_2	x_3	x_4	\dots	x_{n-1}	x_n
-------	-------	-------	-------	---------	-----------	-------

Distillate $\sum x_i, x_i \in \mathbb{F}$




Distillate $\sum x_i, x_i \in \mathbb{F}$

e_1	s_1	x_3	x_4	\dots	x_{n-1}	x_n
-------	-------	-------	-------	---------	-----------	-------

$$\sum_1^n x_i = e_1 + s_1 + \sum_3^n x_i$$

Distillate $\sum x_i, x_i \in \mathbb{F}$



Distillate $\sum x_i, x_i \in \mathbb{F}$

e_1	e_2	s_2	x_4	\dots	x_{n-1}	x_n
-------	-------	-------	-------	---------	-----------	-------


$$\sum_1^n x_i = e_1 + e_2 + s_2 + \sum_4^n x_i$$

Distillate $\sum x_i, x_i \in \mathbb{F}$

e_1	e_2	e_3	e_4	\dots	e_{n-2}	s_{n-1}	x_n
-------	-------	-------	-------	---------	-----------	-----------	-------

$$\sum_1^n x_i = \sum_1^{n-2} e_i + s_{n-1} + x_n$$

Distillate $\sum x_i, x_i \in \mathbb{F}$




Distillate $\sum x_i, x_i \in \mathbb{F}$

e_1	e_2	e_3	e_4	\dots	e_{n-2}	e_{n-1}	s_n
-------	-------	-------	-------	---------	-----------	-----------	-------

$$\sum_1^n x_i = \sum_1^{n-1} e_i + s_n = \sum_1^n x_i^{(1)}$$

and one iterate within this new vector $[x^{(1)}] \dots$

Let's distillate!



Theorem (Zhu-Hayes,09)

Terminating for IEEE-754 : Iterate the distillation

$[x] \rightarrow [x^{(1)}] \rightarrow \dots \rightarrow [x^{(k)}] \rightarrow \dots$ converges towards a stable state $[x^*]$ such that $|x_1^*| < |x_2^*| < \dots < |x_n^*|$, $x_i^* \oplus x_{i+1}^* = x_{i+1}^*$ et $\sum x_i = \sum x_i^*$.

Rmk : the convergence proof uses the round-to-even tie breaking rule.


We need a good stopping criteria!

EFT to sum n floating-point numbers

Rien ne se perd, rien ne se crée, tout se transforme (Anaxagore–Lavoisier)

EFT vector transformation: (Ogita-Rump-Oishi, 05)

$$[p_1, p_2, \dots, p_n] \longmapsto [q_2, q_3, \dots, q_n, \pi_n]$$



For $(p_i)_{1 \leq i \leq n} \in \mathbb{F}$, $p_1 + p_2 = \pi_2 + q_2$, $p_2 + p_3 = \pi_3 + q_3, \dots$,

$$\sum_{i=1}^n p_i = \pi_n + \sum_{i=2}^n q_i, \text{ with } \pi_n = \text{fl}\left(\sum_{i=1}^n p_i\right).$$

A priori stopping criteria: go away conditioning!

Sum2, SumK (Ogita-Rump-Oishi,05)




FIG. 4.2. Outline of Algorithm 4.8 for $n = 5$ and $K = 4$.

Theorem (OgRO,05)

For every p_i in \mathbb{F} ,

$$\sum p_i = \sum p'_i = \cdots = \sum p_i^K = \cdots,$$

$$\text{cond}(\sum p_i^K) = \mathcal{O}(\mathbf{u}^K) \times \text{cond}(\sum p_i)$$

$$\text{Accuracy} \lesssim \mathbf{u} + \text{cond} \times \mathbf{u}^K$$

T. OGITA, S. M. RUMP, AND S. OISHI





FIG. 6.1. Test results for Algorithm 5.3 (Dot2), $n = 100$, 1000 samples.

$$\text{Accuracy} \lesssim \mathbf{u} + \text{cond} \times \mathbf{u}^K$$

T. OGITA, S. M. RUMP, AND S. OISHI



More EFT and accurate algorithms

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms

- Some more results

4 More accuracy: conclusion

More accuracy: more results

Other EFT to compensate

- Multiplication: $a \times b = (a \otimes b) + e$
- FMA: $a \times b + c = FMA(a, b, c) + e_1 + e_2$, and
 $FMA(a, b, c) = RN(a \times b + c)$
- Polynomial evaluation

Other compensated algorithms

- Dot product
- Horner and derivative (dsynthetic div.), de Casteljau (Bernstein), Clenshaw (Chebychev)

Sterbenz's lemma: when subtraction is exact

Theorem (Sterbenz, 72)

In a radix- β floating-point system with subnormal numbers, if x and y are finite floating-point numbers such that

$$\frac{y}{2} \leq x \leq 2y,$$

then $x - y$ is a floating-point number.

Comments

- Exact subtraction for the four IEEE-754 rounding-modes
- Exact subtraction vs. catastrophic cancellation?
- Used in the previous EFT sum proofs

EFT: the multiplication case

Theorem (Dekker)

The multiply rounding-error $xy - (x \otimes y)$ is a floating-point number when no overflow occurs.

When one **FMA is available**, next algorithm computes the multiply-EFT:

$$xy = r_1 + r_2.$$

Algorithm (2ProdFMA)

```
 $r_1 = \text{FMA}(x, y, 0); \quad // r_1 = x \otimes y$   
 $r_2 = \text{FMA}(x, y, -r_1); \quad // r_2 = xy - r_1$ 
```


The multiplication case when no available FMA

Let $C = \beta^s + 1$ where $s < p$, the next algo splits the p -digit floating-point number x in two parts x_h, x_l of $p - s$ and s digits.

Algorithm (split - Veltkamp, 68)

$$r_shift = C \otimes x$$

$$head = x \ominus r_shift$$

$$x_h = r_shift \oplus head$$

$$x_l = x \ominus x_h$$

Theorem (Dekker, 72; Boldo, 06)

When no overflow occurs,

$$x = x_h + x_l.$$

The multiplication case when no available FMA

When no overflow occurs, next algorithm computes the multiply-EFT:

$$xy = r_1 + r_2.$$

Algorithm (2prod - Dekker, 72)

```
( $x_h, x_l$ ) = split( $x$ );  
( $y_h, y_l$ ) = split( $y$ );  
 $r_1 = x \otimes y$ ;  
 $high = -r_1 \oplus (x_h \otimes y_h)$ ;  
 $mid_1 = high \oplus (x_h \otimes y_l)$ ;  
 $mid_2 = mid_1 \oplus (x_l \otimes y_h)$ ;  
 $r_2 = mid_2 \oplus (x_l \otimes y_l)$ ;    //low part
```

The compensated dot product (Ogita-Rump-Oishi, 05)

$$\text{EFT-DotProd: } \sum_{i=1}^n x_i y_i = p + \sum_{i=1}^n \pi_i + \sum_{i=1}^{n-1} \sigma_i = \sum_{i=1}^{2n} z_i.$$

Algorithm (Compensated DotProd)

```
(s1, c1) = 2Prod(x1, y1);  
for i = 2 : n do  
    (pi, πi) = 2Prod(xi, yi);  
    (si, σi) = 2Sum(pi, si-1);  
    ci = ci-1 ⊕ (πi ⊕ σi);  
end for  
return sn ⊕ cn;
```

Theorem (Ogita-Rump-Oishi, 05)

Compensated dot product accuracy $\leq \mathbf{u} \left| \sum_{i=1}^n x_i y_i \right| + \mathcal{O}(n^2 \mathbf{u}^2) \sum |x_i y_i|.$

The compensated dot product (Ogita-Rump-Oishi, 05)

T. OGITA, S. M. RUMP, AND S. OISHI





FIG. 6.1. Test results for Algorithm 5.3 (Dot2), $n = 100$, 1000 samples.

The compensated dot product (Ogita-Rump-Oishi, 05)

T. OGITA, S. M. RUMP, AND S. OISHI



The Horner polynomial evaluation (Graillat-Langlois-Louvet, 07)

EFT-Horner: $\sum_{i=0}^n a_i x^i = p + \sum_{i=0}^{n-1} \pi_i x^i + \sum_{i=0}^{n-1} \sigma_i y^i = p + \pi(x) + \sigma(x).$


Algorithm (Compensated Horner)

```
 $r_n = a_n;$   
for  $i = n - 1 : (-1) : 0$  do  
   $(p_i, \pi_i) = 2Prod(r_{i+1}, x);$   
   $(r_i, \sigma_i) = 2Sum(p_i, a_i);$   
   $q_i = \pi_i \oplus \sigma_i;$  //correcting polynomial coefficient  
end for;  
 $r = r_0 \oplus Horner(q, x);$   
return  $r;$ 
```


Theorem (Langlois-Louvet, 07)

Compensated Horner accuracy $\leq \mathbf{u}|p(x)| + \mathcal{O}(4n^2 \mathbf{u}^2) \sum_{i=0}^n |a_i| |x^i|.$

The compensated Horner evaluation (Louvet's PhD, 07)



The compensated Horner evaluation: FMA variations



Other compensated polynomial evaluations

Recent extensions

- derivative Horner evaluation — synthetic division (Jiang et al., 12)
- Clenshaw algorithm for Chebychev basis (Jiang et al., 11)
- de Casteljau algorithm for Bernstein basis (Jiang et al., 10)

To conclude this Part 1

1 Summing n floating-point numbers: basic steps

- Some algorithms
- Some accuracy bounds

2 Computing sums more accurately

- Some famous old and magic algorithms
- Compensation
- Distillate to understand and go further

3 More EFT and accurate algorithms

- Some more results

4 More accuracy: conclusion

More accurate sums: conclusion

An “ultimately accurate” floating-point summation algorithm cannot be very simple.
J.M. Muller et al. [HFPA,09]

- A lot of algorithms!
- Condition dependency and accuracy bounds for classic ones
- Arbitrarily accurate compensated ones: Sum2, SumK
- Faithfully or correctly rounded recent ones: AccSum, FastAccSum, iFastSum, HybridSum, OnLineExact
- “In place” solutions vs. exponent manipulations

How to choose? Running time performance!

References I



D. H. Bailey.

Twelve ways to fool the masses when giving performance results on parallel computers.
Supercomputing Review, pages 54–55, Aug. 1991.



J. W. Demmel and H. D. Nguyen.

Fast reproducible floating-point summation.

In *Proc. 21th IEEE Symposium on Computer Arithmetic*. Austin, Texas, USA, 2013.



B. Goossens, P. Langlois, D. Parello, and E. Petit.

PerPI: A tool to measure instruction level parallelism.

In K. Jónasson, editor, *Applied Parallel and Scientific Computing - 10th International Conference, PARA 2010, Reykjavík, Iceland, June 6-9, 2010, Revised Selected Papers, Part I*, volume 7133 of *Lecture Notes in Computer Science*, pages 270–281. Springer, 2012.



Y. He and C. Ding.

Using accurate arithmetics to improve numerical reproducibility and stability in parallel applications.

J. Supercomput., 18:259–277, 2001.

References II



J.-M. Hervouet.

Hydrodynamics of free surface flows: Modelling with the finite element method.

John Wiley & Sons, 2007.



N. J. Higham.

Accuracy and Stability of Numerical Algorithms.

SIAM, 2nd edition, 2002.



H. Jiang, R. Barrio, H. Li, X. Liao, L. Cheng, and F. Su.

Accurate evaluation of a polynomial in chebyshev form.

Applied Mathematics and Computation, 217(23):9702 – 9716, 2011.



H. Jiang, S. Graillat, C. Hu, S. Li, X. Liao, L. Cheng, and F. Su.

Accurate evaluation of the k-th derivative of a polynomial and its application.

J. Comput. Appl. Math., 243:28–47, May 2013.



H. Jiang, S. Li, L. Cheng, and F. Su.

Accurate evaluation of a polynomial and its derivative in bernstein form.

Computers & Mathematics with Applications, 60(3):744 – 755, 2010.

References III



P. Langlois and N. Louvet.

How to ensure a faithful polynomial evaluation with the compensated Horner algorithm?

In P. Kornerup and J.-M. Muller, editors, *18th IEEE International Symposium on Computer Arithmetic*, number ISBN 0-7695-2854-6, pages 141–149. IEEE Computer Society, June 2007.



P. Langlois, R. Nheili, and C. Denis.

Numerical Reproducibility: Feasibility Issues.

In *NTMS'2015: 7th IFIP International Conference on New Technologies, Mobility and Security*, pages 1–5, Paris, France, July 2015. IEEE, IEEE COMSOC & IFIP TC6.5 WG.



P. Langlois, R. Nheili, and C. Denis.

Recovering numerical reproducibility in hydrodynamic simulations.

In J. H. P. Montuschi, M. Schulte, S. Oberman, and N. Revol, editors, *23rd IEEE International Symposium on Computer Arithmetic*, number ISBN 978-1-5090-1615-0, pages 63–70. IEEE Computer Society, July 2016.

(Silicon Valley, USA. July 10-13 2016).

References IV



N. Louvet.

Algorithmes compensés en arithmétique flottante : précision, validation, performances.

PhD thesis, Université de Perpignan Via Domitia, Nov. 2007.



J.-M. Muller, N. Brisebarre, F. de Dinechin, C.-P. Jeannerod, V. Lefèvre, G. Melquiond, N. Revol, D. Stehlé, and S. Torres.

Handbook of Floating-Point Arithmetic.

Birkhäuser Boston, 2010.



T. Ogita, S. M. Rump, and S. Oishi.

Accurate sum and dot product.

SIAM J. Sci. Comput., 26(6):1955–1988, 2005.



Open TELEMAT-MASCARET. v.7.0, Release notes.

www.opentelemat.org, 2014.



R. W. Robey, J. M. Robey, and R. Aulwes.

In search of numerical consistency in parallel programming.

Parallel Comput., 37(4-5):217–229, 2011.

References V



S. M. Rump.

Ultimately fast accurate summation.

SIAM J. Sci. Comput., 31(5):3466–3502, 2009.



S. M. Rump, T. Ogita, and S. Oishi.

Accurate floating-point summation – part I: Faithful rounding.

SIAM J. Sci. Comput., 31(1):189–224, 2008.



M. Taufer, O. Padron, P. Saponaro, and S. Patel.

Improving numerical reproducibility and stability in large-scale numerical simulations on gpus.

In *IPDPS*, pages 1–9. IEEE, 2010.



O. Villa, D. G. Chavarría-Miranda, V. Gurumoorthi, A. Márquez, and S. Krishnamoorthy.

Effects of floating-point non-associativity on numerical computations on massively multithreaded systems.

In *CUG 2009 Proceedings*, pages 1–11, 2009.

References VI



V. Weaver and J. Dongarra.

Can hardware performance counters produce expected, deterministic results?

In *3rd Workshop on Functionality of Hardware Performance Monitoring, 2010*, pages 1–11, Atlanta, USA, 2010.



D. Zaparanuks, M. Jovic, and M. Hauswirth.

Accuracy of performance counter measurements.

In *IEEE International Symposium on Performance Analysis of Systems and Software, ISPASS 2009, April 26–28, 2009, Boston, Massachusetts, USA*, pages 23–32, 2009.



Y.-K. Zhu and W. B. Hayes.

Correct rounding and hybrid approach to exact floating-point summation.

SIAM J. Sci. Comput., 31(4):2981–3001, 2009.



Y.-K. Zhu and W. B. Hayes.

Algorithm 908: Online exact summation of floating-point streams.

ACM Trans. Math. Software, 37(3):37:1–37:13, Sept. 2010.